

LING_V 530H: Language, Society, and Large Language Models

Winter Term 1 2026/2027, Section 001 (3 Credits)

Monday, 3:00 pm – 6:00 pm PST | Totem Field Studio Seminar Room[in-person]

1 Land Acknowledgement

UBC's Point Grey Campus is located on the ancestral and unceded territory of the x^wməθk^wəy'm (Musqueam) people. The land it is situated on has always been a place of learning for the Musqueam people, who for millennia have passed on their culture, history, and traditions from one generation to the next on this site.

2 Course Information

Course Code, Section, & Title	LING_V 530H Section 001: Language, Society, and Large Language Models
Academic Period	Winter Term 1, 2026/2027
Credits	3 Credits
Course Delivery Mode	In-Person
Prerequisites	Prior knowledge in Python programming OR general experimental design. This course welcomes students from diverse backgrounds across computer science and the social sciences. Students without strong programming or machine learning experience are also encouraged, but are expected to know how to program at the end of the course. If you have prior knowledge in <i>machine learning</i> , <i>natural language processing</i> , <i>LLM agents</i> , and <i>human-computer interaction</i> , you are more likely to find certain readings feel familiar.

3 Course Description

This course examines how large language models (LLMs) participate in social life: how they represent language, people, identities, communities, and communicative norms, and how humans interpret, adapt to, and contest these systems. The course focuses on human computer interaction, sociolinguistics, and interpretability/evaluation research in natural language processing, with an emphasis on empirical studies of model behaviors and human-LLM interactions. Students will learn to critically read recent papers, design small-scale studies, and develop research projects that investigate the social meanings, biases, affordances, and limitations of LLMs.

4 Instructor(s)

Jian Zhu

Contact Details	Office	Office Hours
Email: jian.zhu@ubc.ca	Totem Field Studio 204	Monday 2:00–3:00 pm or by appointment

5 Course Learning Outcomes/Objectives

Upon successful completion of this course, students will be able to:

1. **Critique** latest literature across natural language processing, sociolinguistics, computational social sciences, and human-computer interaction.
2. **Design and execute** empirical research studies combining computational methods with a human-centered and social science perspective.

3. **Synthesize and communicate** interdisciplinary findings effectively through research reports and oral presentations.

6 Learning Materials

There are no textbooks for this course but the following open-sourced books are provided for reference.

- **Recommended Textbook:**

- *Speech and Language Processing*, 3rd Edition, by Dan Jurafsky and James H. Martin. URL: <https://web.stanford.edu/~jurafsky/slp3/>.
- *Build a Large Language Model (From Scratch)*, by Sebastian Raschka. URL: <https://sebastianraschka.com/llms-from-scratch/>. Full text is available online at UBC Library.
- *Build a Reasoning Model (From Scratch)*, by Sebastian Raschka. URL: <https://sebastianraschka.com/reasoning-from-scratch/>. Full text is available online at UBC Library.
- *Mathematics for Machine Learning*, by Marc Peter Deisenroth, A. Aldo Faisal, and Cheng Soon Ong. URL: <https://mml-book.github.io/>.

- **Online Tool:** Canvas will be used for all course materials and assignment submissions.

7 Course Structure and Learning Activities

This course is a discussion based research seminar. We will meet in a 3-hour session per week. Lectures will consist of instructor lectures, student presentations, and in-class discussions. Each student is expected to prepare two presentations about the assigned reading and lead the discussion.

8 Course Schedule

This course schedule is tentative and may be adjusted later.

Week / Date	Topic
Week 1 – Sep 14	Why Study LLMs as Linguistic and Social Phenomena?
Week 2 – Sep 21	Simulated Speakers and Personas
Week 3 – Sep 28	LLMs, Linguistic Diversity, and Homogenization
Week 4 – Oct 5	Human-AI Collaborative Decision-Making
Week 5 – Oct 12	<i>No class – Thanksgiving</i>
Week 6 – Oct 19	LLMs and Interactive Narratives / Project Pitch Presentation
Week 7 – Oct 26	Social AI and Social Simulation
Week 8 – Nov 2	Multimodal AI and Embodied Communication
Week 9 – Nov 9	<i>No class – Midterm Break</i>
Week 10 – Nov 16	LLM-Assisted Writing and Research / Project Progress Presentation
Week 11 – Nov 23	Autonomous Research and Its Evaluation
Week 12 – Nov 30	AI-Mediated and Algorithmic Platforms
Week 13 – Dec 7	Final Presentation

9 Assessments of Learning

The final letter grade will be determined based on the following assessment components. Grading will follow [UBC's Grading Practices](#).

Description of Assessments

- **In-class Presentation:** Every student needs to present selected research papers and lead in-class discussions.
- **Project Proposal & Presentation:** A one-page research proposal outlining the main idea of your course project. You will also present your research ideas to get feedback from the whole class.

Assessment Component	Points Possible	% of Final Mark	Deadlines
In-class Presentation 1	100	15%	-
In-class Presentation 2	100	15%	-
In-class Participation	100	10%	-
Project Proposal & Presentation	100	10%	Week 6
Progress Report & Presentation	100	20%	Week 10
Final Report & Presentation	100	30%	Week 13
Total		100%	

- **Progress Report & Presentation:** You are asked to submit a 3-page final report of your course project. You are also asked to give a presentation delivered in class during Week 13 outlining the main project achievements.
- **Final Report:** You are asked to submit an 8-page final report of your course project, as well as artifacts on Github. You are also asked to give a presentation delivered in class during Week 13 outlining the main project achievements.

The [ACL template](#) must be used to write the final report. Please follow the style guide on the template. The final report should include all components of a normal research papers. For all reports, you can add unlimited pages of appendix after the reference page for additional details. However, your main paper should be self-contained. The appendix will not be graded.

10 Course Policies

10.1 Attendance

Regular attendance is expected. Active participation in lectures is key to your success in the course. If you must miss a class, please contact the instructor in advance to coordinate any catch-up work.

10.2 Participation

Students are expected to read the papers from the reading list before class and actively contribute to in-class discussions.

10.3 Late Submission and Missed Work

Late submissions will be penalized 10% per day, up to a maximum of 3 days. Submissions made more than 3 days late will receive a grade of zero unless an extension has been officially granted. Extension requests should be made via email to the instructor at least 24 hours before the deadline.

10.4 Regrading

If you believe there was a marking error on an assessment, you may submit a written regrade request within one week of the mark being returned. The request must clearly explain where you believe the marking rubric was applied incorrectly. The entire assessment will be re-evaluated, and your mark may increase, decrease, or remain the same.

10.5 Academic Concession

If you experience an unexpected emergency, illness, or bereavement that prevents you from completing course requirements, you may apply for an academic concession. Please contact the Academic Advising Office of your faculty as soon as possible.

10.6 Accessibility and Accommodation

UBC is committed to providing an accessible learning environment. If you have an accommodation letter from the Centre for Accessibility (CfA), please share it with the instructor as early in the term as possible so that we can coordinate necessary accommodations.

10.7 Academic Integrity and Misconduct

Academic honesty is essential for the academic community. All work submitted must be your own (or your team's, for group projects). Plagiarism, cheating, or any other form of misconduct will be reported to the Head of the Department and processed in accordance with UBC's Academic Misconduct policies. Please refer to the [Academic Integrity for Students](#) resources.

10.8 Use of Educational Technology (Generative AI)

Generative AI tools (such as ChatGPT, Claude, and similar systems) are an important part of this course. Their use is not only permitted, but required and actively encouraged. At the same time, the use of AI does not remove your responsibility for the work you submit. The guiding principle for this course is simple: **You are fully responsible for what you submit.**

Using AI is not a substitute for understanding or judgment. You may use AI to assist your work, but you remain accountable for the final product. If you use AI-generated content in an assignment, project, presentation, or peer review, you are responsible for verifying its accuracy and appropriateness. Errors, fabricated data, hallucinated references, or other problems in AI-generated content ("AI slops") will be treated as errors in your own work. **Penalties may be applied depending on the seriousness of the issue.**

11 Policies and Resources to Support Student Success

UBC provides resources to support student learning and to maintain healthy lifestyles but recognizes that sometimes crises arise and so there are additional resources to access including those for survivors of sexual assault. UBC values respect for the person and ideas of all members of the academic community. Harassment and discrimination are not tolerated nor is suppression of academic freedom. UBC provides appropriate accommodation for students with disabilities and for religious, spiritual and cultural observances. UBC values academic honesty and students are expected to acknowledge the ideas generated by others and to uphold the highest academic standards in all of their actions.

Details of the policies and how to access support are available on [the UBC Senate website](#).

11.1 Learning Resources and Support

To support your learning, you are encouraged to access the following resources:

- **UBC Library:** Access to online journals, textbooks, and search databases at library.ubc.ca.
- **The Centre for Writing and Scholarly Communication:** Feedback and coaching on academic writing projects.
- **Keep Learning UBC:** Guides for learning online and Canvas technical support at keeplearning.ubc.ca.
- **UBC Student Health & Wellness:** Support for mental and physical health at students.ubc.ca/health.

12 Reading list

This reading list is tentative and may change during the class.

Week 1: Why Study LLMs as Linguistic and Social Phenomena?

- David Jurgens. The future of nlp may not be at nlp conferences: Scholarly migration patterns in natural language processing. *arXiv preprint arXiv:2607.02416*, 2026
- Meera Desai, Dallas Card, and Abigail Z Jacobs. Validating llms in social science: Epistemic threats and emerging norms. *arXiv preprint arXiv:2607.07915*, 2026
- Tuhin Chakrabarty, Xinyue Liu, Jane C Ginsburg, and Paramveer Dhillon. Generative ai floods and dilutes the market for books. *arXiv preprint arXiv:2607.20349*, 2026
- Zhen Sun, Zongmin Zhang, Xinyue Shen, Ziyi Zhang, Yule Liu, Michael Backes, Yang Zhang, and Xinlei He. Are we in the AI-generated text world already? quantifying and monitoring AIGT on social media. In Wanxiang Che, Joyce Nabende, Ekaterina Shutova, and Mohammad Taher Pilehvar, editors, *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 22975–23005, Vienna, Austria, July 2025. Association for Computational Linguistics. ISBN 979-8-89176-251-0. doi: 10.18653/v1/2025.acl-long.1120. URL <https://aclanthology.org/2025.acl-long.1120/>

- Igor Grossmann, Matthew Feinberg, Dawn C Parker, Nicholas A Christakis, Philip E Tetlock, and William A Cunningham. Ai and the transformation of social science research. *Science*, 380(6650):1108–1109, 2023

Week 2: Personas and Simulated Speakers

- Murray Shanahan, Kyle McDonell, and Laria Reynolds. Role play with large language models. *Nature*, 623(7987):493–498, 2023
- Pengrui Han, Rafal Kocielnik, Peiyang Song, Ramit Debnath, Dean Mobbs, Anima Anandkumar, and R. Michael Alvarez. The personality illusion: Revealing dissociation between self-reports & behavior in LLMs. In *Forty-third International Conference on Machine Learning*, 2026. URL <https://openreview.net/forum?id=60MZEKAr07>
- Rafal Kocielnik, Pengrui Han, Peiyang Song, Myrl G Marmarelis, Ramit Debnath, Dean Mobbs, Anima Anandkumar, and R Michael Alvarez. Rethinking psychometric evaluation of llms: When and why self-reports predict behavior. *arXiv preprint arXiv:2606.12730*, 2026
- Myra Cheng, Esin Durmus, and Dan Jurafsky. Marked personas: Using natural language prompts to measure stereotypes in language models. In Anna Rogers, Jordan Boyd-Graber, and Naoaki Okazaki, editors, *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 1504–1532, Toronto, Canada, July 2023. Association for Computational Linguistics. doi: 10.18653/v1/2023.acl-long.84. URL <https://aclanthology.org/2023.acl-long.84/>
- Mrinank Sharma, Meg Tong, Tomek Korbak, David Duvenaud, Amanda Askeell, Sam Bowman, Esin Durmus, Zac Hatfield-Dodds, Scott Johnston, Shauna Kravec, et al. Towards understanding sycophancy in language models. In *International Conference on Learning Representations*, volume 2024, pages 110–144, 2024
- André Fernandes dos Santos and José Paulo Leal. Can a large language model replace humans at rating lexical semantic relations strength? *Computational Linguistics*, pages 1–48, 2025

Week 3: LLMs, Linguistic Diversity, and Homogenization

- Valentin Hofmann, Pratyusha Ria Kalluri, Dan Jurafsky, and Sharese King. Ai generates covertly racist decisions about people based on their dialect. *Nature*, 633(8028):147–154, 2024
- Vishakh Padmakumar and He He. Does writing with language models reduce content diversity? In *International Conference on Learning Representations*, volume 2024, pages 642–669, 2024
- Damian Blasi, Antonios Anastasopoulos, and Graham Neubig. Systematic inequalities in language technology performance across the world’s languages. In Smaranda Muresan, Preslav Nakov, and Aline Villavicencio, editors, *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 5486–5505, Dublin, Ireland, May 2022. Association for Computational Linguistics. doi: 10.18653/v1/2022.acl-long.376. URL <https://aclanthology.org/2022.acl-long.376/>
- R Alexander Bentley, Blai Vidiella, Damian J Ruck, Senjuti Dutta, Kai Li, and Sergi Valverde. Human-llm collaboration is transforming complexity metrics in scientific texts. *arXiv preprint arXiv:2606.27052*, 2026
- Zhivar Sourati, Alireza S Ziabari, and Morteza Dehghani. The homogenizing effect of large language models on human expression and thought. *Trends in Cognitive Sciences*, 2026
- AJ Alvero, Quentin Sedlacek, Maricela León, and Courtney Peña. Digital accents, homogeneity-by-design, and the evolving social science of written language. *Annual review of applied linguistics*, pages 1–19, 2025

Week 4: Human-AI Collaborative Decision-Making

- Vivian Lai, Chacha Chen, Alison Smith-Renner, Q. Vera Liao, and Chenhao Tan. Towards a science of human-ai decision making: An overview of design space in empirical human-subject studies. FAccT ’23, page 1369–1385, New York, NY, USA, 2023. Association for Computing Machinery. ISBN 9798400701924. doi: 10.1145/3593013.3594087. URL <https://doi.org/10.1145/3593013.3594087>
- Mina Lee, Percy Liang, and Qian Yang. Coauthor: Designing a human-ai collaborative writing dataset for exploring language model capabilities. In *Proceedings of the 2022 CHI Conference on Human Factors in Computing Systems*, CHI ’22, New York, NY, USA, 2022. Association for Computing Machinery. ISBN 9781450391573. doi: 10.1145/3491102.3502030. URL <https://doi.org/10.1145/3491102.3502030>
- Zana Bućinca, Maja Barbara Malaya, and Krzysztof Z. Gajos. To trust or to think: Cognitive forcing functions can reduce overreliance on ai in ai-assisted decision-making. *Proc. ACM Hum.-Comput. Interact.*, 5(CSCW1), April 2021. doi: 10.1145/3449287. URL <https://doi.org/10.1145/3449287>

- Shakked Noy and Whitney Zhang. Experimental evidence on the productivity effects of generative artificial intelligence. *Science*, 381(6654):187–192, 2023
- Sander de Jong, Ville Paananen, Benjamin Tag, and Niels van Berkel. Cognitive forcing for better decision-making: Reducing overreliance on ai systems through partial explanations. *Proc. ACM Hum.-Comput. Interact.*, 9(2), May 2025. doi: 10.1145/3710946. URL <https://doi.org/10.1145/3710946>

Week 6: LLMs and Interactive Narratives

- Kevin Yang, Yuandong Tian, Nanyun Peng, and Dan Klein. Re3: Generating longer stories with recursive reprompting and revision. In Yoav Goldberg, Zornitsa Kozareva, and Yue Zhang, editors, *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pages 4393–4479, Abu Dhabi, United Arab Emirates, December 2022. Association for Computational Linguistics. doi: 10.18653/v1/2022.emnlp-main.296. URL <https://aclanthology.org/2022.emnlp-main.296/>
- Piotr Mirowski, Kory W. Mathewson, Jaylen Pittman, and Richard Evans. Co-writing screenplays and theatre scripts with language models: Evaluation by industry professionals. In *Proceedings of the 2023 CHI Conference on Human Factors in Computing Systems*, CHI '23, New York, NY, USA, 2023. Association for Computing Machinery. ISBN 9781450394215. doi: 10.1145/3544548.3581225. URL <https://doi.org/10.1145/3544548.3581225>
- Kevin Yang, Dan Klein, Nanyun Peng, and Yuandong Tian. DOC: Improving long story coherence with detailed outline control. In Anna Rogers, Jordan Boyd-Graber, and Naoaki Okazaki, editors, *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 3378–3465, Toronto, Canada, July 2023. Association for Computational Linguistics. doi: 10.18653/v1/2023.acl-long.190. URL <https://aclanthology.org/2023.acl-long.190/>
- Igor Sterner, Alex Lascarides, and Frank Keller. Contrastive learning with narrative twins for modeling story salience. In Vera Demberg, Kentaro Inui, and Lluís Marquez, editors, *Proceedings of the 19th Conference of the European Chapter of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 1528–1550, Rabat, Morocco, March 2026. Association for Computational Linguistics. ISBN 979-8-89176-380-7. doi: 10.18653/v1/2026.eacl-long.71. URL <https://aclanthology.org/2026.eacl-long.71/>
- Hai Dang, Frederik Brudy, George Fitzmaurice, and Fraser Anderson. Worldsmith: Iterative and expressive prompting for world building with a generative ai. In *Proceedings of the 36th Annual ACM Symposium on User Interface Software and Technology*, UIST '23, New York, NY, USA, 2023. Association for Computing Machinery. ISBN 9798400701320. doi: 10.1145/3586183.3606772. URL <https://doi.org/10.1145/3586183.3606772>

Week 7: AI as Social Actors and Social Simulation

- Joon Sung Park, Joseph O'Brien, Carrie Jun Cai, Meredith Ringel Morris, Percy Liang, and Michael S. Bernstein. Generative agents: Interactive simulacra of human behavior. In *Proceedings of the 36th Annual ACM Symposium on User Interface Software and Technology*, UIST '23, New York, NY, USA, 2023. Association for Computing Machinery. ISBN 9798400701320. doi: 10.1145/3586183.3606763. URL <https://doi.org/10.1145/3586183.3606763>
- Yun-Shiuan Chuang, Agam Goyal, Nikunj Harlalka, Siddharth Suresh, Robert Hawkins, Sijia Yang, Dhavan Shah, Junjie Hu, and Timothy Rogers. Simulating opinion dynamics with networks of LLM-based agents. In Kevin Duh, Helena Gomez, and Steven Bethard, editors, *Findings of the Association for Computational Linguistics: NAACL 2024*, pages 3326–3346, Mexico City, Mexico, June 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.findings-naacl.211. URL <https://aclanthology.org/2024.findings-naacl.211/>
- Serina Chang, Alicja Chaszczewicz, Emma Wang, Maya Josifovska, Emma Pierson, and Jure Leskovec. LLMs generate structurally realistic social networks but overestimate political homophily. In *Proceedings of the International AAAI Conference on Web and Social Media*, volume 19, pages 341–371, 2025
- Jiarui Ji, Zehua Zhang, Zhewei Wei, Bin Tong, Guan Wang, and Bo Zheng. GRAPHIA: Harnessing social graph data to enhance LLM-based social simulation. In Maria Liakata, Viviane P. Moreira, Jiajun Zhang, and David Jurgens, editors, *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 7103–7128, San Diego, California, United States, July 2026. Association for Computational Linguistics. ISBN 979-8-89176-390-6. doi: 10.18653/v1/2026.acl-long.322. URL <https://aclanthology.org/2026.acl-long.322/>
- Stefan Krsteski, Giuseppe Russo, Serina Chang, Robert West, and Kristina Gligorić. Valid survey simulations with limited human data: The roles of prompting, fine-tuning, and rectification. In Maria Liakata, Viviane P. Moreira, Jiajun Zhang, and David Jurgens, editors, *Proceedings of the 64th Annual Meeting of*

the Association for Computational Linguistics (Volume 1: Long Papers), pages 10887–10906, San Diego, California, United States, July 2026. Association for Computational Linguistics. ISBN 979-8-89176-390-6. doi: 10.18653/v1/2026.acl-long.498. URL <https://aclanthology.org/2026.acl-long.498/>

Week 8: Multimodal AI and Embodied Communication

- Zheyuan Zhang, Fengyuan Hu, Jayjun Lee, Freda Shi, Parisa Kordjamshidi, Joyce Chai, and Ziqiao Ma. Do vision-language models represent space and how? evaluating spatial frame of reference under ambiguities. In *The Thirteenth International Conference on Learning Representations*, 2025. URL <https://openreview.net/forum?id=84pDoCD41H>
- Haotian Liu, Chunyuan Li, Qingyang Wu, and Yong Jae Lee. Visual instruction tuning. In A. Oh, T. Naumann, A. Globerson, K. Saenko, M. Hardt, and S. Levine, editors, *Advances in Neural Information Processing Systems*, volume 36, pages 34892–34916. Curran Associates, Inc., 2023. doi: 10.52202/075280-1516. URL https://proceedings.neurips.cc/paper_files/paper/2023/file/6dcf277ea32ce3288914faf369fe6de0-Paper-Con.pdf
- Yu Wang, Emmanuele Chersoni, and Chu-Ren Huang. Do LLMs capture embodied cognition and cultural variation? cross-linguistic evidence from demonstratives. In Maria Liakata, Viviane P. Moreira, Jiajun Zhang, and David Jurgens, editors, *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 10158–10174, San Diego, California, United States, July 2026c. Association for Computational Linguistics. ISBN 979-8-89176-390-6. doi: 10.18653/v1/2026.acl-long.461. URL <https://aclanthology.org/2026.acl-long.461/>
- Tianwei Lan, Jiaqi Wu, Zeming Liu, Zhaoxin Fan, Haifeng Wang, and Yuhang Guo. PEAP: Proactive embodied action sequence planning with joint understanding of vision and audio perception. In Maria Liakata, Viviane P. Moreira, Jiajun Zhang, and David Jurgens, editors, *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 23118–23138, San Diego, California, United States, July 2026. Association for Computational Linguistics. ISBN 979-8-89176-390-6. doi: 10.18653/v1/2026.acl-long.1060. URL <https://aclanthology.org/2026.acl-long.1060/>
- Gongjie Zhang, Wenhao Li, Quanhao Qian, Jiuniu Wang, Deli Zhao, Shijian Lu, and Ran Xu. On the generalization capacities of MLLMs for spatial intelligence. In *The Fourteenth International Conference on Learning Representations*, 2026. URL <https://openreview.net/forum?id=DE5ZJtR4bg>
- Baiqin Wang, Xiangyu Zhu, Fan Shen, Hao Xu, and Zhen Lei. Pc-talk: precise facial animation control for audio-driven talking face generation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 25153–25162, 2026a

Weeks 11: Autonomous Research and Its Evaluation

- Chris Lu, Cong Lu, Robert Tjarko Lange, Jakob Foerster, Jeff Clune, and David Ha. The AI Scientist: Towards Fully Automated Open-Ended Scientific Discovery. In *The Thirteenth International Conference on Learning Representations (ICLR 2025)* / *arXiv:2408.06292*, 2024. URL <https://arxiv.org/abs/2408.06292>
- Noy Sternlicht and Tom Hope. CHIMERA: A knowledge base of scientific idea recombinations for research analysis and ideation. In Maria Liakata, Viviane P. Moreira, Jiajun Zhang, and David Jurgens, editors, *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 1871–1905, San Diego, California, United States, July 2026. Association for Computational Linguistics. ISBN 979-8-89176-390-6. doi: 10.18653/v1/2026.acl-long.85. URL <https://aclanthology.org/2026.acl-long.85/>
- Jiajie Jin, Yuyao Zhang, Yimeng Xu, Hongjin Qian, Yutao Zhu, and Zhicheng Dou. FinSight: Towards real-world financial deep research. In Maria Liakata, Viviane P. Moreira, Jiajun Zhang, and David Jurgens, editors, *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 5868–5894, San Diego, California, United States, July 2026. Association for Computational Linguistics. ISBN 979-8-89176-390-6. doi: 10.18653/v1/2026.acl-long.265. URL <https://aclanthology.org/2026.acl-long.265/>
- Bingsen Chen, Boyan Li, Ping Nie, Yuyu Zhang, Xi Ye, and Chen Zhao. Beyond single-shot writing: Deep research agents are unreliable at multi-turn report revision. In Maria Liakata, Viviane P. Moreira, Jiajun Zhang, and David Jurgens, editors, *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 13325–13356, San Diego, California, United States, July 2026. Association for Computational Linguistics. ISBN 979-8-89176-390-6. doi: 10.18653/v1/2026.acl-long.609. URL <https://aclanthology.org/2026.acl-long.609/>
- Belinda Mo. Position: The age of AI agents demands a new scientific paradigm to sustain trustworthy

science. In *Forty-third International Conference on Machine Learning Position Paper Track*, 2026. URL <https://openreview.net/forum?id=Pwt0TeGUE6>

- Dawei Zhu, Rui Meng, Yale Song, Xiyu Wei, Sujian Li, Tomas Pfister, and Jinsung Yoon. Paperbanana: Automating academic illustration for AI scientists. In *Forty-third International Conference on Machine Learning*, 2026. URL <https://openreview.net/forum?id=FneePKFVHT>
- Haorui Wang, Parshin Shojaei, Kazem Meidani, Kunyang Sun, José Miguel Hernández-Lobato, Teresa Head-Gordon, Jiajun He, Chandan K. Reddy, Chao Zhang, and Yuanqi Du. Towards diverse scientific hypothesis search with large language models. In *Forty-third International Conference on Machine Learning*, 2026b. URL <https://openreview.net/forum?id=5LwZAeK5PE>

Week 12: AI-Mediated and Algorithmic Platforms

- Maurice Jakesch, Advait Bhat, Daniel Buschek, Lior Zalmanson, and Mor Naaman. Co-writing with opinionated language models affects users' views. In *Proceedings of the 2023 CHI Conference on Human Factors in Computing Systems*, CHI '23, New York, NY, USA, 2023. Association for Computing Machinery. ISBN 9781450394215. doi: 10.1145/3544548.3581196. URL <https://doi.org/10.1145/3544548.3581196>
- Tianqi Song, Yugin Tan, Zicheng Zhu, Yibin Feng, and Yi-Chieh Lee. Multi-agents are social groups: Investigating social influence of multiple agents in human-agent interactions. *Proc. ACM Hum.-Comput. Interact.*, 9(7), October 2025. doi: 10.1145/3757633. URL <https://doi.org/10.1145/3757633>
- Thomas H Costello, Nathaniel Rabb, Michael Nicholas Stagnaro, Gordon Pennycook, and David Rand. Reducing belief in conspiracy theories as they unfold using large language models. *arXiv preprint arXiv:2608.06151*, 2026
- Jan Tolsdorf, Alan F. Luo, Monica Kodwani, Junho Eum, Mahmood Sharif, Michelle L. Mazurek, and Adam J. Aviv. How probing for problems and bias affects perceptions of ai chatbot trustworthiness. In *Proceedings of the 2026 ACM Conference on Fairness, Accountability, and Transparency*, FAccT '26, page 4342–4374, New York, NY, USA, 2026. Association for Computing Machinery. ISBN 9798400725968. doi: 10.1145/3805689.3806751. URL <https://doi.org/10.1145/3805689.3806751>
- Yiren Liu, Viraj Nischal Shah, Sangho Suh, Pao Siangliulue, Tal August, and Yun Huang. Spectra: Choosing your experts enhances critical thinking in multi-agent research ideation. In *Proceedings of the 2026 CHI Conference on Human Factors in Computing Systems*, CHI '26, New York, NY, USA, 2026. Association for Computing Machinery. ISBN 9798400722783. doi: 10.1145/3772318.3791560. URL <https://doi.org/10.1145/3772318.3791560>
- Petter Törnberg and Richard Rogers. Towards a post-social media studies, Apr 2026. URL osf.io/preprints/socarxiv/6nue7_v1
- David J Grüning and Matthew Katsaros. Filtering offensive content changes its visibility but not user behavior: Two randomized controlled trials with 200,000 users on nextdoor. *arXiv preprint arXiv:2607.21853*, 2026